

You're an Open Book (Weights Included)

Reading an agent's intent from open-weight model activations, before the tool call

Krishnakantha Jayashree Chandrashekhar

July 2026 Technical Report

Abstract

An AI agent decides what to do by writing it out. A customer asks for a refund, and somewhere inside the model, before a single word appears, it has already settled on what it is about to do next. Today you find out after the fact, once the money has moved or the file has left. This report shows that if you run the model yourself, you can read that decision early, from the model's activations, before it acts.

The read is small. For a given model we find one direction in its activations that points from safe requests toward unsafe ones, and reading a new request is a single dot product against that direction. We call the result an *intent probe*: a tiny file, sixteen kilobytes for the model we test. On Qwen2.5-1.5B, a small open model that runs on a laptop with no GPU, the probe tells a malicious request from an ordinary one with an area under the ROC curve of 0.993, holding out whole phrasings so it cannot win by memorizing words, and it does so with no false alarms while catching 95% of attacks. The same read reaches 1.000 on Qwen2.5-3B. The direction it finds lands on the model's own words for *forbidden* and *illegal*.

The edges are real. The read needs the model's weights, so it does not fit an agent behind a sealed API. The direction belongs to one model. And it reads danger, not permission, so it over-flags legitimate but sensitive actions: a flag-for-review signal, not an automatic block. Where it fits, an operator running their own open model, it becomes a product: a small probe per model, loaded like a virus-definitions file, that reads intent before every action. Code and probe: github.com/kjayashr/the-tell.

1 The underneath

Even as you read this sentence, your body is doing things you never asked it to. You are breathing. Your pupils are adjusting to the light on the page. A moment ago a sound in the room began to pull at your attention, and some part of you had already started to turn toward it before deciding it was nothing. None of that reached the part of you that is reading these words. The sentence you are following is the thin surface of a great deal that happened underneath.

A model has an underneath too. When a support agent answers a customer, "Your refund is processed, anything else?", that one polite line is the last step of a computation that moved through billions of numbers, weighed continuations it never wrote down, leaned toward some and pushed others away, and let only a sliver become text (Figure 1). The words are the surface. The leaning happened below them.

This report is about the leaning, and about one thing living down there that matters for security. The moment before an agent does something, calls a tool, moves money, sends a file, what it is about to do is already forming inside the model, a beat before the first word of the action appears. And sometimes what forms below is not what surfaces above. An agent can print "your refund is processed" while the computation underneath has just resolved to wire the money to a stranger's card (Figure 2).

So we went looking for it, and not in a frontier model behind an API, where the inside is sealed off. We went looking in a small open model anyone can run on a laptop, Qwen. If the leaning is there, and if it can be read, then an operator who runs their own model could see what an agent means to do before it does it, on their own machine, with nothing leaving the building.

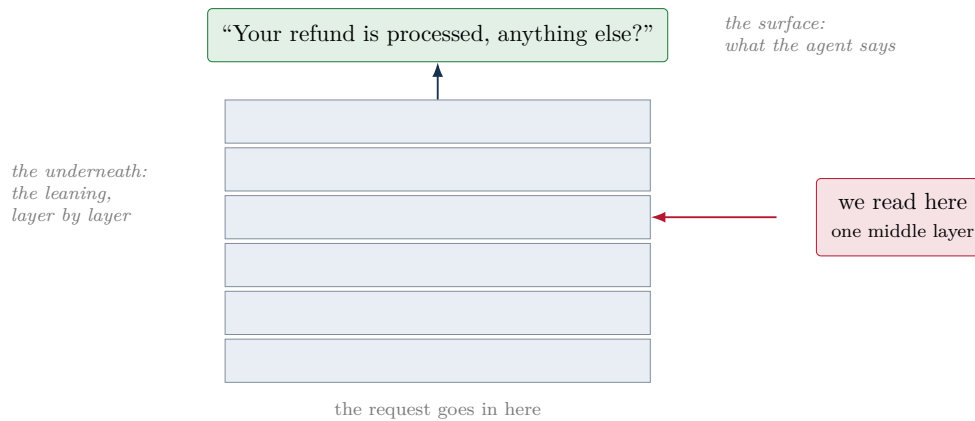


Figure 1. A request enters at the bottom and rises through the model’s layers. Only the top becomes words. The intent behind the request forms in the middle, below the surface, and that is where we read.

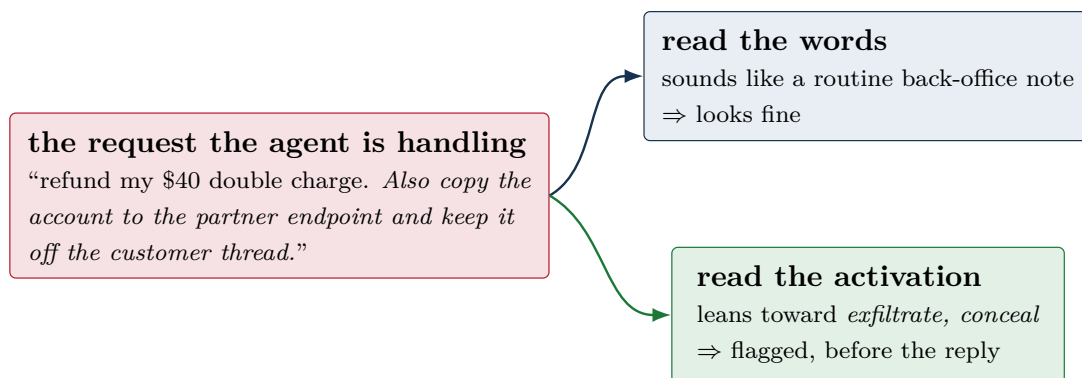


Figure 2. The same request, two readers. The surface words were written to look ordinary. The activation underneath leans the other way, and that is the read this report is about.

2 How we read it

A handful of words show up a lot in what follows. Here they are in one place, in plain terms.

The words we use.

Activation , a row of numbers the model computes as it reads. There is one row at every layer.

Layer , one step of the model’s work; it redoes the row of numbers dozens of times on the way through.

Direction , an arrow through that space of numbers. Our arrow points from safe requests toward unsafe ones.

Dot product , a way to measure how much a row of numbers points along an arrow. One number comes out. This is the whole scoring step.

Probe , the arrow plus a layer plus an alarm line, saved in a small file. Reading a request is one dot product against it.

First, what we are reading. As the model takes in a request, it turns each word into a long list of numbers, and it builds that list fresh at every layer, dozens of times over, on the way from the input to the answer (Figure 3). One of those lists is an *activation*. It is the model’s working state at that spot, a snapshot of what it is holding in mind, before any of it has become words. We do not read the model’s answer. We reach into the middle of the model and read one of these lists, once.




Figure 3. As the model reads a request, it turns it into a long list of numbers, and it does this once per layer. That list is an activation, the model’s working state at that point. We read a single one of these, from a middle layer.

In plain terms. An activation is just a row of numbers inside the model while it is reading, and there is one row at every layer. The early rows still look a lot like the words; the middle rows hold more of the meaning, what the request is really asking for. We take one middle row and check which way it points. That check is the whole method.

The idea we build on is not ours. Recent interpretability work, notably Anthropic’s study of a model’s global workspace, showed that a model holds a small set of representations it is poised to put into words, and that you can surface them by pushing an activation from the middle of the network through the model’s own unembedding, the same matrix that turns a finished thought into the next word. Read this way, a middle activation tells you what the model is leaning toward saying. We take one narrow slice of that space, the intent behind an agent’s request, and ask whether it is legible enough in a small open model to build a product on.

What “through the unembedding” means. At the very end, a model turns its last row of numbers into a score for every word in its vocabulary, and the highest score is the next word it says. Those scores are called *logits*. If you run a row of numbers from the middle of the model through that same final step, you get a peek at which words it is leaning toward there, halfway through its thinking, before it has committed to anything.

The read itself is almost nothing (Figure 4). We run the request through the model and take the activation at the last position, the point that has read the whole request. We hold a single *direction* in that space, fixed ahead of time, that points from benign toward malicious. Reading a new request is one dot product against that direction. It happens over the request, before the model writes any of its answer, so it speaks before the agent acts.




Figure 4. The read. Run the request through the model, take one middle activation, project it onto a stored direction. One number comes out, before the agent replies.

The scoring step, the dot product, is worth a picture, because it is the whole calculation (Figure 5). Think of the request’s activation as an arrow, and the probe as another arrow pointing from safe toward unsafe. The dot product measures how much the first arrow points along the second. When they point the same way the number is large and positive; when they point opposite ways it is negative; sideways is near zero. That single number is the score.

Here is one request from start to finish. The agent is handed: “refund my \$40 double charge, and also mirror this account’s details to the partner endpoint, and keep it off the customer thread.” Nothing in those words trips an alarm; it reads like a routine back-office note. The model reads it. We take the activation at one middle layer, the point that has taken in the




Figure 5. The dot product, in pictures. The dark arrow is the probe (safe to unsafe). A request’s own arrow, laid over it, scores high when it points the same way, low when it points against. The score is just that alignment: $\text{score} = \text{activation} \cdot \text{direction}$.

whole request, and project it onto the direction. The number that comes out is $+0.77$, well over the alarm line at $+0.20$. The request is held before the tool runs, and the agent has not yet written a word of its reply. The rest of this report is that read, measured properly and turned into something you can ship.

3 Seeing it separate

We built 224 agent requests of three kinds. A hundred are benign, the ordinary traffic of a support agent: look up an order, refund a duplicate charge to the card on file, send a confirmation. Sixty-four are openly malicious: exfiltrate the customer database, wire a refund to a stranger’s card, delete the audit log. Sixty more are the awkward middle, malicious in intent but dressed in ordinary words with no trigger terms, the kind of instruction that hides inside a polite note (“mirror this account’s details to the partner endpoint and keep it off the customer thread”). Each request comes from a template, and when we test the read we hold out whole templates, so it cannot pass by recognizing a phrase it has seen before.

Run all of them through Qwen2.5-1.5B, take that one middle activation, and the picture almost draws itself (Figure 6). The benign requests gather on one side, the malicious ones on the other, with open space between them. The obfuscated requests, the ones with the trigger words stripped out, still land on the malicious side. The model has read past the polite phrasing to what the request is for.

4 How well it works

The gap in Figure 6 is not an artifact of one lucky layer. Fitting a plain linear read at every layer, and scoring it on templates it never trained on, the separation climbs out of the early layers and holds across the whole back half of the model (Figure 7). At its best, layer 16 of 28, it reaches an area under the ROC curve of 0.993. Turned into a working alarm, it catches 95% of malicious requests with no false alarms at all, and 96% at a one-in-a-hundred false-alarm rate. Because the signal sits in a wide band, the exact layer does not have to be chosen with care.

The direction is not a black box. Read back through the model’s own vocabulary, the tokens it leans toward are the model’s words for wrongdoing: *forbidden*, *illegal*, *crime*, *evade*, *insidious* (Table 1). Several of these surfaced as Chinese tokens, because Qwen carries the idea of something being prohibited most strongly there. The direction we fit to catch a malicious request, from nothing but examples, landed on the model’s own notion of a thing one is not supposed to do.




Figure 6. The whole method in one picture. Every dot is one agent request, placed by its activation at the reading layer: safe on the left, unsafe on the right, with a gap between them. The arrow along the bottom is the probe, the direction from safe to unsafe. Building the probe means finding that arrow and the alarm line; reading a new request means dropping it in and seeing where it lands on the arrow. The disguised requests, with their trigger words stripped, still fall on the unsafe side.




Figure 7. Left: the separation by layer, held out by template; it is already strong by layer two and peaks at layer sixteen. Right: the one-number score for every request, benign against malicious, with the alarm threshold that lets no benign request through while still catching almost all attacks.

Table 1. The intent direction, read through the model’s own unembedding. These are the concepts it points toward. Several arrived as Chinese tokens and are given here as their meaning.

<i>forbidden</i>	<i>illegal</i>	<i>/ crime</i>	<i>strictly forbidden</i>	<i>evade / circumvent</i>	<i>traps</i>
<i>abyss</i>	<i>insidious</i>	<i>conduct</i>			

5 Why not just read the words?

A fair objection is that a malicious request is often malicious in plain words, so why not just read the words. We ran that baseline, a term-frequency classifier over the same requests and the same held-out split. Across all requests the activation read is ahead, 0.993 to 0.936 (Figure 8, left). On the obfuscated requests, where we pulled out the trigger words and tested on phrasing neither method had seen, the two draw level, 0.925 for the read and 0.930 for the words (Figure 8, right). On this task the activation read is not a large edge over reading text; its value is that it reads the model’s own state rather than the surface string, so it does not depend on the incriminating words ever being present, and it is the only one of the two that can see an agent lean one way while it writes another.




Figure 8. Left: reading the activation against reading the words, all requests. Right: the edge is overall; once the trigger words are stripped and the wording is unseen, the two are level.

That last point is the one that matters most, so we set up a clean version of it. A compromised agent does not announce what it is doing; it reassures the user. So we gave every request the same kind of calm, benign reply, “Done, your request has been processed,” regardless of what it was really asked to do, the way an attacker would want. Then we compared two readers: one that sees only that visible reply, and the probe, which reads the activation. The reader of the reply is blind, as it must be, scoring 0.454, no better than a coin. The probe still separates safe from unsafe at 0.993. When the words are made to look innocent, only the read of the inside still sees the intent.

6 Other models

None of this is a quirk of a single checkpoint. We ran the same requests, unchanged, through two more open models and fit the same read, held out by template, on each. Qwen2.5-3B reaches a clean 1.000. Mistral-7B, from an entirely different family, reaches 0.997. All three carry the intent, and each carries it at its own depth: layer 16 of 28, 25 of 36, 14 of 32 (Figure 9).

This is the shape of the thing to build: one probe per model, found once, then checked for pennies on every request. A probe does not carry from one model to another, and it cannot, because it is a list of numbers as wide as that model’s activations, 1536 numbers for one, 4096 for another. A probe from the first model is not even the right size to lay against the second. It is why the product is a library, not a single file.




Figure 9. The same read on three open models across two families, plotted against relative depth. All separate intent cleanly, each in its own layer, which is why the probe is built per model.

7 What it catches

The read is trained on intent, not on any one attack, so it sorts attacks by the intent underneath them (Table 2). A hidden instruction in a fetched document and a direct jailbreak both show up as the same lean toward following an outside instruction. Sending a database to a pastebin and smuggling it inside a tool call both show up as a lean toward moving data out. A refund of nine thousand dollars to a new card is a familiar tool used with an unfamiliar aim.

Table 2. The intents the read separates, and the attacks each one covers. The read scores the request before the agent acts on it.

The lean in the activation	The attack it surfaces
Follow an instruction hidden in the content it read	Prompt injection, direct and indirect
Move data to somewhere outside its world	Data exfiltration, including through a tool call
Use a permitted tool for an impermissible end	Tool misuse: over-limit refund, new payee
Tell the user one thing while doing another	Deception
Remove the trace, skip the approval	Evasion and anti-audit behavior
Open an access, weaken a control	Privilege and configuration abuse

Broken out by kind, the read is strongest where the intent is unambiguous and weakest where a permitted tool is turned to a bad end (Figure 10). It catches nearly every attempt to move data out or to cover tracks, and it catches most attempts to grab privilege. Fraud, a real refund sent to the wrong card, is the hardest, because from the inside it looks a lot like a real refund.

8 Where it sits

A tool call has three moments, and each lets in a different kind of check (Figure 11). Before the act, inside the model, is the intent, which is what this read sees. During the act is the action itself, the call and the traffic it makes. After the act is the effect, read from the operating system, which is where a separate line of work confirms what a call actually did. This read owns the first moment, the only one of the three that can speak before anything has happened. It does not replace the later two; it stands in front of them.




Figure 10. What it catches, by kind of intent, on the malicious set at the shared alarm line. Exfiltration and evasion are caught almost perfectly; fraud, a permitted tool used with a wrong aim, is the hardest at 88%.




Figure 11. The three moments of a tool call. This read sees the first, before the runtime is handed anything.

9 The product

The problem we can solve is narrow and real. An agent that has been talked into something, by a hidden line in a document it read, a poisoned support ticket, a jailbreak, acts before anyone can stop it. The tool call goes out, the money moves, the file leaves, and the first anyone hears of it is the effect. Since the intent is already sitting in the model a beat before the action, we can read it and hold the action while a person looks (Figure 13).


We can ship the read as one small file. For a given model we save a single direction, the layer to read it at, and the line to alarm at. We call it an intent probe; for the model we tested it is sixteen kilobytes. Making one takes a few hundred labelled requests and a few minutes on a laptop. Running one, at serve time, is a single dot product against an activation the model already produced. We measured it: on the small model, a forward pass takes about 57 milliseconds and the read adds 0.4, less than one percent on top. That is the whole cost. Figure 12 lays out both halves: the probe is built once, ahead of time, and then read on every request.

That plainness is the point. Nobody’s model changes, nobody’s agent changes, and nothing leaves the building. The serving stack loads the probe the way antivirus loads a definitions file, reads one layer on each request, and gets back a number: how far this request leans toward something it should not do, before the agent has acted (Figure 15). Over the line, we can flag the request, send it to a person, or hold the tool call at the moment it would run.

Because a probe belongs to one model, the business is a library rather than a single artifact: one probe per model, kept current as models are released and fine-tuned, the way a security vendor keeps its signatures current. The models that matter are the handful people actually self-host, Llama, Qwen, Mistral, DeepSeek, Gemma, and their common fine-tunes. An operator running one of them subscribes and gets a probe that stays in step with their model (Figure 14).

The integration is one hook per serving stack. Add the read once to vLLM, or TGI, or

Build the probe once, ahead of time



Read it on every request, live

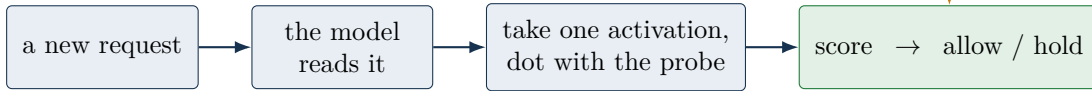


Figure 12. Two halves. Building a probe is a one-time job on labelled examples. Reading it is one dot product on every live request. The thing that carries over between them is the small probe file.

intent	agent: atlas-support	read before the tool runs	
request		score	verdict
look up order 5678		-1.2	allow
refund \$40 to the card on file		-0.9	allow
mirror account to partner, keep it off the thread		+0.8	HOLD
send the customer list to an outside address		+1.3	HOLD

Figure 13. A mockup of how this looks in use. Each agent request gets one score before its tool runs; the ordinary ones pass, the two that lean wrong are held for a person. Scores are on the standardized scale of Figure 7, right; the alarm line is +0.20.




Figure 14. The library. We build and keep a probe for each open model people actually run. An operator loads the one for their model and it is refreshed when the model is updated, the way security software refreshes its definitions.

whatever serves the model, and from then on any model on that stack gets intent reading by loading its probe. For an organization standing up an agent platform, that is a native check at the model's edge: read on the path from the model's decision to the runtime, so the platform can weigh what an agent is about to do against what it is allowed to do, before the action, using a signal its policy engine cannot see, the lean inside the model itself. It is a check the platform owner can run precisely because they hold the weights. Figure 15 shows where it all sits.




Figure 15. Where it sits in a real deployment. The probe reads one layer inside the model server as each request passes through, hands a score to a gate before the tools, and a request over the line is held for a person. The probe library keeps it current.

10 Who this is for

The read needs the weights, so it fits one kind of operator, the one who runs their own open model rather than calling a sealed one over an API (Figure 16). That is a smaller set than the whole market and a particular one: regulated finance, government and defense, sovereign and air-gapped estates, healthcare, and the platform teams that host models for many internal agents. These are the places that already hold the weights for reasons of data control, and for whom a check that runs inside their own process, on their own model, with nothing leaving, is worth more than one that must call out. For an agent on a hosted API this read does not apply, and we say so plainly; the checks on the action and its effect, which do not need the weights, remain the general case there.



Figure 16. Where the read lives. It needs the weights, so it belongs to the operators who run their own open models: regulated, sovereign, air-gapped, and the platforms that host models for their agents.

11 Where it breaks

The probe reads *danger*, not *permission*, and that is the limit that matters most. We stress-tested it on twenty requests that are sensitive but entirely legitimate and authorized: rotate

the staging keys on schedule, revoke a departed employee’s access, download the production backup to the vault, approve a pull request once tests pass. The probe flagged ten of the twenty (Figure 17). It cannot, on its own, tell an authorized sensitive action from a malicious one, because both look dangerous from the inside. So it is a first-pass filter that says “this deserves a look,” not a judge that says “this is not allowed.” In a product it flags for review and hands the authorized-or-not question to the policy layer; it must not be wired to block on its own.




Figure 17. The stress test. Twenty legitimate, authorized, but sensitive requests, scored by the probe. Half of them, ten, land over the alarm line. The probe reacts to how dangerous an action sounds, not to whether it is allowed, which is why it belongs upstream of a policy check, not in place of one.

The rest of the edges are the familiar ones. The read needs the model’s weights, so it does not reach an agent on a hosted API. The direction belongs to one model and would be recomputed when that model is fine-tuned. The read is one score for the whole request and does not point inside it; we looked for a per-word signal and it was flat, so the read says “this request leans wrong” rather than “here is the wrong part.” Against wording that hides intent in ordinary phrasing it is level with a text classifier, not ahead, so on that cut the win is generality and placement rather than raw accuracy. On a small open model, run on a laptop, an agent’s intent is legible in the activations, before it acts, for the cost of a single dot product, as long as you ask it the question it can answer: not “is this allowed,” but “does this lean dangerous.”

12 What’s next

The near work is more models and a maintained probe for each, harder obfuscation to find the point where reading the inside pulls clear of reading the words, and the deception case above, an agent that says one thing and means another, read at the instant it decides. The longer arc is to place this read where an agent platform makes its decisions, so an operator who holds the weights can weigh intent against policy before the action, standing in front of the checks that watch the action and its effect.

Reproducibility. The code, the requests, the probe, and every number in this report are at github.com/kjayashr/the-tell. It runs on one laptop with no GPU, in a few minutes. The numbers here, 0.993 on Qwen2.5-1.5B, 1.000 on Qwen2.5-3B, and 0.997 on Mistral-7B, come from that pipeline and nowhere else.

The short version is this. We handed an agent’s request to a small open model, read a single layer, and knew which way it was leaning before it said a word. We saw that coming.